

Lab 4. Running Mapreduce

Exercise 1. Install mapreduce

Initially we will do everything on node 3.

First we install the relevant packages.

```
yum install -y hadoop-mapreduce hadoop-doc
```

Verify that the sample jar file now works and shows the example programs.

```
yarn jar /usr/share/doc/hadoop-0.20-mapreduce/examples/hadoop-examples.jar
```

Replace any existing content in the file `/etc/hadoop/conf/mapred-site.xml` with the following:

```
<?xml version="1.0" encoding="UTF-8"?>
<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>
```

Next, we will configure YARN to load an auxiliary service that can perform the shuffle phase of the mapreduce algorithm by adding the following to the existing `yarn-site.xml` config:

```
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.aux-services.mapreduce_shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
```

Now we will configure hadoop application managers to store their temporary files within user directories in HDFS by adding the following lines to the existing yarn-site.xml config file.

```
<property>
<name>yarn.app.mapreduce.am.staging-dir</name>
<value>/user</value>
</property>
```

Make sure that this temporary directory exists in hdfs.

```
su -c 'hdfs dfs -mkdir /tmp' hdfs
su -c 'hdfs dfs -chmod -R 1777 /tmp' hdfs
```

Next restart :YARN to allow the changes to take effect.

```
service hadoop-yarn-resourcemanager restart
service hadoop-yarn-nodemanager restart
ps -u yarn w
```

Now let's create some sample data for our mapreduce job.

```
find /etc -exec strings {} \+ > /tmp/etc.words
find /usr/share/doc -exec strings {} \+ > /tmp/doc.words
ls -lh /tmp*.words
```

Let's upload this file to hdfs.

```
su - student
hdfs dfs -put /tmp/etc.words /upload/
hdfs dfs -put /tmp/doc.words /upload/
```

Now we'll run the mapreduce job, placing the output in the users home directory.

```
Yarn jar /usr/lib/hadoop-0.20-mapreduce/hadoop-examples.jar
wordcount /upload/etc.words
```

Look at the temporary directories that were created by the mapreduce job.

```
hdfs dfs -ls /home/<username>
```

Look at the yarn logs.

```
su -c 'ls -R /yarn/logs/application_*0001' root
```

Now let's get the results from the mapreduce job.

```
hdfs dfs -get /home/<username>/result1  
ls -l result1/
```

Let's look, for example, for all words that occurred between 500 and 600 times.

```
awk '{if ($2 > 500 && $2 < 600) {print}}' result1/part-r-00000
```

Now we'll add the node managers on the additional nodes in our cluster.

First, let's see how long it takes to run the job with only one node.

```
time yarn jar /usr/lib/hadoop-0.20-mapreduce/hadoop-examples.jar  
wordcount \ /upload/doc.words result2
```

Now let's push our MapReduce/YARN configs to the other nodes.

```
cd /etc/hadoop.conf  
for n in 1 2 4; do scp mapred-site.xml yarn-site.xml  
node$n:/etc/hadoop/conf; done
```

Start the nodemanager daemon on the other nodes.

```
allnodes service hadoop-yarn-nodemanager start  
yarn node list | grep Nodes
```

Now, as the student user run the mapreduce job again and time it to see what the difference is.

```
time yarn jar /usr/lib/hadoop-0.20-mapreduce/hadoop-examples.jar  
wordcount /upload/doc.words result3
```

Now run the same job with multiple reducers.

```
yarn jar /usr/lib/hadoop-0.20-mapreduce/hadoop-examples.jar \  
wordcount mapreduce.job.reduces=3 /upload/doc.words results_multi  
  
hdfs dfs -ls results_multi  
hdfs dfs -get results_multi  
grep ^watch part-r-0000"*
```

Note that the data has been partitioned across multiple files. We can merge them back by using standard Linux commands.

```
sort part-r-0000* -o merged  
grep ^watch merged
```

However we can do the same thing more easily by using the `-getmerge` option to `hdfs dfs`

```
hdfs dfs -getmerge results_multi merged
```